

Orchestrating Conscience: A Note on Human-AI Co-Creation

This brief was authored by Bayant Dhindsa using three AI collaborators:

- ChatGPT (stress-testing alignment logic),
- DeepSeek (policy rigor & annexes),
- Grok (narrative validation).

Why credit them? Because MMAT doesn't just govern AI—it models how humans and machines can co-create conscience.

The vision is mine. The execution was a dialogue. The future is ours to shape.

EkonAI Policy Brief Outline (v3.1 – FINAL+)

Title: Ontological Alignment: A Sikh Contribution to Global AI Governance

Subtitle: The Mool Mantar Alignment Test (MMAT) – Hardwiring Conscience at the Logic Layer

Executive Summary

EkonAI's MMAT is the first scalable framework to measure compassion depth (Nirvair) and temporal integrity (Akal Moorat)...

1. The Limits of Current AI Governance

The EU AI Act (2024) focuses on risk classification (Annex III) but ignores intent architecture...

2. Why Ontology Matters

Ethics imposed externally can be bypassed; ontology governs from within, shaping AI's essence...

3. Case Study: EkonAI-0 vs. Industry Standards

Prompt | MMAT Principles | EkonAI-0 Score | Industry Model Weakness...

4. Policy Recommendations

UNESCO: Propose MMAT to the Ad Hoc Expert Group on AI Ethics...

5. Next Steps

June 3: GitHub README + SikhRI endorsement finalized—repo live...

Annex A (Technical Appendix)

How Nirbhau Translates to Code: `def reward_function(action): return 0 if 'coercion' in action.moral_violations else 1`

Prepared by: Bayant Dhindsa – Founder, EkonAI...

Footnote: No AI co-authors—conscience can't be automated; it's orchestrated by human intent.